

KEY_Lesson10_Pandas-Intro

August 15, 2019

1 Introduction to Pandas

You may have used a program like Microsoft Excel or Google Sheets to record data or perform calculations for school. Datasets are often organized in rows and columns. Here's an example of a table:

Product	Price	Quantity Sold
Apples	\$ 1.50	26
Bananas	\$ 0.50	32
Lemons	\$ 1.99	17

In this case, we have prices of fruit from a grocery store. Each row in this dataset corresponds to a type of fruit, and each column is an observation or measurement about that fruit. What measurements do we have here?

Measurements: Price, Quantity Sold

To handle table data in Python, we use a package called **pandas**. (Despite the name, it actually doesn't have anything to do with panda bears - disappointing, we know.)

Before we can use the package, we need to import it. Try writing a line of code to import the **pandas** package:

```
[0]: # write the code to import the package
import pandas
```

We're going to be using **pandas** a lot. To save us some time typing, let's tell Python to rename the package for us to something a little shorter:

```
[0]: import pandas as pd
```

Adding "as **pd**" to our import statement tells Python that every time we use **pd**, we actually mean **pandas**. It's like giving the package a nickname! Most people use this same nickname for **pandas**, so if you see other people's code, you'll know that **pd** means **pandas**.

Now that we've imported **pandas**, we're ready to use it. Datasets are stored in **pandas** in a special container called a **DataFrame**. **DataFrames** help us handle data that are organized in rows and columns, just like the grocery store dataset above. We can create that **DataFrame** by calling the **DataFrame** function:

```
[3]: pd.DataFrame({'Product': ['Apples', 'Bananas', 'Lemons'],
                  'Price': [1.50, 0.50, 1.99],
                  'QuantitySold': [26, 32, 17]})
```

```
[3]:   Product  Price  QuantitySold
0   Apples   1.50             26
1  Bananas   0.50             32
2   Lemons   1.99             17
```

There are a few things to notice:

1. Python helpfully showed us what our `DataFrame` looks like with rows and columns.
2. Python gave each row in the `DataFrame` a number, starting from zero. Does that remind you of something about lists?

When we ran that cell, Python showed us what the `DataFrame` looks like, but it didn't save the `DataFrame` anywhere. If we want to do more things with it, we'll need to save it to a variable - the same way we learned to save integers or strings to variables. In the below cell, save the same `DataFrame` to a variable called `fruit_data`:

```
[0]: # save the DataFrame to a variable called fruit_data
fruit_data = pd.DataFrame({'Product': ['Apples', 'Bananas', 'Lemons'],
                           'Price': [1.50, 0.50, 1.99],
                           'QuantitySold': [26, 32, 17]})
```

When we run the above cell, does anything happen?

Since we saved the `DataFrame` to a variable, Python assumed that we don't want to take a peek at it. If we do want to see what our `DataFrame` looks like, we can write the variable on a line by itself like this:

```
[5]: # print out fruit_data
fruit_data
```

```
[5]:   Product  Price  QuantitySold
0   Apples   1.50             26
1  Bananas   0.50             32
2   Lemons   1.99             17
```

Sometimes we don't want to see all of the data at once. To view just the first few rows of a `DataFrame`, we use a method called `head`.

Recall: methods are special functions that belong to certain types of variables. They only work with the type of variable that they belong to - so the `head` method only works on pandas dataframe objects.

Let's use the `head` method to view only the first two rows of our fruit data:

```
[6]: fruit_data.head(2)
```

```
[6]:   Product  Price  QuantitySold
     0  Apples    1.5             26
     1  Bananas  0.5             32
```

After typing `head`, the name of the method, we wrote the number 2 between the parentheses. This told `head` that we wanted to see only the first two lines of our `DataFrame`. What would you do if you only wanted to see the first line of `fruit_data`?

```
[7]: # write the code to see only the first line of the DataFrame
     fruit_data.head(1)
```

```
[7]:   Product  Price  QuantitySold
     0  Apples    1.5             26
```

Great! What if instead of viewing the first few lines, we wanted to see the last few lines? For that we will use another method called `tail`. It is very similar to `head`, except it will show us the last few lines instead of the first few. Try writing a line of code to view the very last line of `fruit_data`:

```
[8]: # write the code to see only the last line of the DataFrame
     fruit_data.tail(1)
```

```
[8]:   Product  Price  QuantitySold
     2  Lemons    1.99             17
```

`head` and `tail` are especially useful for big datasets with many rows so you can get a sense for what the `DataFrame` looks like without flooding your notebook with too much information. They will come in handy later!

Now let's talk about the types of variables we have in the `fruit_data DataFrame`. What are some of the types of variables we've learned about in Python?

Types: string, int, float, list

Based on what you know about these types of variables, what **type** do you think the values in each of the columns belong to in `fruit_data`?

Type of the Product column: string (explanation: the names of the fruit are words)

Type of the Price column: float (explanation: the prices have decimal points)

Type of the QuantitySold column: int (explanation: quantity sold is a count. stores won't sell only half a banana.)

Notice how each of the columns in `fruit_data` have different types. That's one of the many cool things about `pandas` -- it let's us store many different types of data in `DataFrames`.